SAVILE ROW Manual 1.6.4

Peter Nightingale

1 Introduction

SAVILE ROW is a constraint modelling tool. It provides a high-level language for the user to specify their constraint problem, and automatically translates that language to the input language of a constraint solver. SAVILE ROW is a research tool, designed to enable research into reformulation of constraint models, therefore it is designed for flexibility. It is easy to add new transformation rules and develop new translation pipelines.

This manual covers the basics of using SAVILE ROW. It does not cover adding new rules or translation pipelines. SAVILE ROW reads the ESSENCE′ modelling language, which is described in the ESSENCE′ description. SAVILE ROW converts constraint models expressed in ESSENCE′ into the solver input format, in a process that has some similarities to compiling a high-level programming language. Like a compiler, SAVILE ROW applies some optimisations to the model (for example, constant folding and common subexpression elimination).

1.1 Problem Classes and Instances

The distinction between *problem classes* and *problem instances* will be important in this manual. It is easiest to start with an example. Sudoku in general is a problem class, and a particular Sudoku puzzle is an instance of that class. The problem class is described by writing down the rules of Sudoku, i.e. that we have a 9×9 grid and that each row, column and subsquare must contain all values 1..9. A particular instance is described by taking the class and adding the clues, i.e. filling in some of the squares. The set of clues is a *parameter* of the class – in this case the only parameter. Adding the parameters to a class to make a problem instance is called *instantiation*.

In typical use, SAVILE ROW will read a problem class file and also a parameter file, both written in ESSENCE′. It will instantiate the problem class and unroll all quantifiers and matrix comprehensions, then perform reformulations and flattening of nested expressions before producing output for a constraint solver.

1.2 Translation Pipelines and Solver Backends

The current version of SAVILE ROW has two translation pipelines. The first works on problem instances only. Every parameter of the problem class must be given a value in the parameter file. Output produced by this pipeline may be entirely flattened or it may still contain some nested expressions depending on the output language.

- Minion – Output is produced in the Minion 3 language for Minion 1.8 or later. The model produced is not entirely flat, it makes use of nested `watched-or` and `watched-and` constraints.

- Gecode – Output is produced for Gecode in the entirely flat, instance-level language Flatzinc for use by the `fzn-gecode` tool.

- SAT – Output is produced in DIMACS format for use with any SAT solver.

- Minizinc – Output is produced in an instance-level, almost flat subset of the Minizinc modelling language. The Minizinc output follows the Minion output as closely as possible. The `mzn2fzn` tool can then be used to translate the problem instance for use with a number of different solvers.

The second translation pipeline does not instantiate the model, instead it performs class-level reformulation and flattening. This pipeline is experimental and incomplete. There is one backend that produces Dominion Input Language for the Dominion solver synthesiser.

MINION is well supported by SAVILE ROW so we will use MINION as the reference in this document. MINION is a fast scalable constraint solver. However, modelling problems directly in MINION's input language is time-consuming and tedious because of its primitive structure (it can be compared to writing a complex program in assembly language).

2   Installing and Running SAVILE ROW

SAVILE ROW is distributed as an archive with the following contents:

- The Java source code (in `src`) licensed with GPL 3.

- The compiled classes in a JAR file named `savilerow.jar`.

- Executable scripts `savilerow` and `savilerow.bat` for running SAVILE ROW.

- This document and the ESSENCE′ description, in `doc`.

- Various example ESSENCE′ files and parameter files, in `examples`.

- Required Java libraries in `lib`

- The Minion solver 1.8 in `bin`

Three archives are provided for Linux, Mac and Windows. These are largely the same, with the main difference being the Minion executable.

A recent version of Java is required on all platforms. The included JAR file was compiled with Oracle Java 1.7.

2.1   Running SAVILE ROW on Linux and Mac

Download the appropriate archive and unpack it somewhere convenient. Open a terminal and navigate to the SAVILE ROW directory. Use the script named `savilerow`. One of the simplest ways of running SAVILE ROW is given below.

```
./savilerow problemFile.eprime parameterFile.param
```

The first argument is the problem class file. This is a plain text file containing the constraint problem, expressed in the ESSENCE′ language.

The second argument (`parameterFile.param`) is the parameter file (again in the ESSENCE′ language). This contains the data for the problem instance. The parameter file can be omitted if the problem has no parameters (no `given` statements in the preamble).

2.2   Running SAVILE ROW on Windows

Download the appropriate archive and unpack it somewhere convenient. Start `cmd.exe` and navigate to the SAVILE ROW folder. Use the script named `savilerow.bat`. Command line options are identical to the Linux and Mac version. One of the simplest ways of running SAVILE ROW is the following:

```
savilerow.bat problemFile.eprime parameterFile.param
```

2.3   Solution Files

For the Minion, SAT and Gecode backends, SAVILE ROW is able to run the solver and parse the solution (or set of solutions) produced by the solver. These solutions are translated back into ESSENCE′. For each `find` statement in the model file (i.e. each statement that declares decision variables), the solution file contains a matching `letting` statement. For example, if the model file contains the following `find` statement:

```
find M : matrix indexed by [int(1..2), int(1..2)] of int(1..5)
```

The solution file could contain the following `letting` statement.

```
letting M = [ [2,3 ; int(1..2)],
              [1,2 ; int(1..2)]
              ; int(1..2) ]
```

2.4   File Names

The input files for SAVILE ROW typically have the extensions `.eprime` and `.param`. These extensions allow SAVILE ROW to identify the model and parameter file on the command line. If these files have a different extension (or no extension) then SAVILE ROW must be called in a slightly different way:

```
./savilerow -in-eprime problemFile -in-param parameterFile
```

Given input file names `<problemFile>` and `<parameterFile>`, output files have the following names by default.

- For Minion output, `<parameterFile>.minion` (or if there is no parameter file, `<problemFile>.minion`).

- For Gecode output, `<parameterFile>.fzn` (or if there is no parameter file, `<problemFile>.fzn`).

- For SAT output, `<parameterFile>.dimacs` (or if there is no parameter file, `<problemFile>.dimacs`).

- For Minizinc output, `<parameterFile>.mzn` (or if there is no parameter file, `<problemFile>.mzn`).

- For Dominion output, there is typically no parameter file so the output file is named after the problem file. If the problem file ends with `.eprime`, then the extension is removed. The extension `.dominion` is added.

When SAVILE ROW parses a single solution from the output of a solver, it produces a file named `<parameterFile>.solution` (or if there is no parameter file, `<problemFile>.solution`). When there are multiple solutions (e.g. when using the `-all-solutions` flag) the solution files are numbered (for example, nurses.param.solution.000001 to nurses.param.solution.000871).

If SAVILE ROW runs Minion or a SAT solver it produces a file `<parameterFile>.info` (or if there is no parameter file, `<problemFile>.info`) containing solver statistics.

Finally, a file named `<parameterFile>.aux` (or `<problemFile>.aux`) is also created. This contains the symbol table and is read if SAVILE ROW is called a second time to parse a solution.

2.5   Command Line Options

2.5.1   Mode

SAVILE ROW has two modes of operation, Normal and ReadSolution. Normal is the default, and in this mode SAVILE ROW reads an ESSENCE′ model file and optional parameter file and produces output for some solver. In some cases it will also run a solver and parse the solution(s), producing ESSENCE′ solution files.

In ReadSolution mode, SAVILE ROW reads a Minion solution table file. The solution table file is created by running Minion with its `-solsout` flag. The solution or solutions saved by Minion are converted to ESSENCE′ solution files. In the process, any decision variables that were removed by SAVILE ROW are restored, and all auxiliary variables that were introduced by SAVILE ROW are removed.

The mode is specified as follows.

```
-mode [Normal | ReadSolution]
```

2.5.2 Specifying input files

The options `-in-eprime` and `-in-param` specify the input files as in the example below.

```
savilerow -in-eprime sonet.eprime -in-param sonet1.param
```

These flags may be omitted if the filenames end with `.eprime` and `.param` respectively.

The option `-params` may be used to specify the parameters on the command line. For example, suppose the `nurse.eprime` model has two parameters. We can specify them on the command line as follows. The format of the parameter string is identical to the format of a parameter file (where, incidentally, the `language` line is optional).

```
savilerow -in-eprime nurse.eprime -params "letting n_nurses=4 letting Demand=[[1,0,1,0],[0,2,1,0]]"
```

2.5.3 Specifying output format

In Normal mode, there are five output formats as described in Section 1.2. The output format is specified using one of the following five command-line options.

```
-minion
-gecode
-sat
-minizinc
-dominion
```

The output filename may be specified as follows. In each case there is a default filename so the flag is optional. Default filenames are described in Section 2.4.

```
-out-minion <filename>
-out-gecode <filename>
-out-sat <filename>
-out-minizinc <filename>
-out-dominion <filename>
```

In addition, the file names for solution files, solver statistics files and aux files may be specified as follows. Once again there are default filenames described in Section 2.4.

```
-out-solution <filename>
-out-info <filename>
-out-aux <filename>
```

2.5.4 Optimisation Levels

The optimisation levels (`-O0` to `-O3`) provide an easy way to control how much optimisation SAVILE ROW does, without having to switch on or off individual optimisations. The default is `-O2`, which is intended to provide a generally recommended set of optimisations. The rightmost `-O` flag on the command line is the one that takes precedence.

```
-O0
```

The lowest optimisation level, `-O0`, turns off all optional optimisations. SAVILE ROW will still simplify expressions (including constraints). Any expression containing only constants will be replaced with its value. Some expressions have quite sophisticated simplifiers that will run even at `-O0`. For example, `allDiff([x+y+z, z+y+x, p, q])` would simplify to `false` because the first two expressions in the `allDiff` are symbolically equal after normalisation. `-O0` will do no common subexpression elimination, will not unify equal variables, and will not filter the domains of variables.

`-O1`

   `-O1` does optimisations that are very efficient in both space and time. Variables that are equal are unified, and a form of common subexpression elimination is applied (Active CSE, described below).

`-O2`

   In addition to the optimisations performed by `-O1`, `-O2` performs filtering of variable domains and aggregation (both of which are described in the following section). `-O2` is the default optimisation level.

`-O3`

   In addition to the optimisations performed by `-O2`, `-O3` enables associative-commutative common subexpression elimination (described below).

2.5.5 Translation Options

Common Subexpression Elimination

SAVILE ROW currently implements four types of common subexpression elimination (CSE). Identical CSE finds and eliminates syntactically identical expressions. This is the simplest form of CSE, however it can be an effective optimisation. Active CSE performs some reformulations (for example applying De Morgan's laws) to reveal expressions that are semantically equivalent but not syntactically identical. Active CSE subsumes Identical CSE. Active CSE is enabled by default as part of `-O2`. Associative-Commutative CSE (AC-CSE) works on the associative and commutative (AC) operators $+$, $*$, $/\backslash$ (and) and $\backslash/$ (or). It is able to rearrange the AC expressions to reveal common subexpressions among them. AC-CSE is not enabled by default. It would normally be used in conjunction with Identical or Active CSE. Finally, Active AC-CSE combines one active reformulation (integer negation) with AC-CSE, so for example Active AC-CSE is able to extract $y - z$ from the three expressions $x + y - z$, $w - x - y + z$, and $10 - y + z$, even though the sub-expression occurs as $y - z$ in one and $z - y$ in the other two. Active AC-CSE is identical to AC-CSE for And, Or and Product, it differs only on sum.
   The following flags control CSE. The first, `-no-cse`, turns off all CSE. The other flags each turn on one type of CSE.

```
-no-cse
-identical-cse
-active-cse
-ac-cse
-active-ac-cse
```

Variable Deletion

SAVILE ROW can remove a decision variable (either variables declared with `find` or auxiliary variables introduced during tailoring) when the variable is equal to a constant, or equal to another variable. This is often a useful optimisation. It can be enabled using the following flag.

```
-deletevars
```

Domain Filtering

It can be useful to filter the domains of variables. In SAVILE ROW this is done by running the translation pipeline to completion and producing a Minion file, then running Minion usually with the options `-preprocess SACBounds -outputCompressedDomains`. With these options Minion performs conventional propagation plus SAC on the lower and upper bound of each variable (also known as shaving). The filtered domains are then read back in to SAVILE ROW. The translation process is started again at the beginning. Thus domain filtering can benefit the entire translation process: variables can be removed (with `-deletevars`), constraint expressions can be simplified or even removed

(if they are entailed), the number of auxiliary variables may be reduced. In some cases domain filtering can enable another optimisation, for example on the BIBD problem it enables associative-commutative CSE to do some very effective reformulation.

Domain filtering can be used with Minion, Gecode, SAT and Minizinc output (but Minion is always used to perform the domain filtering regardless of the output solver). It is switched on using the following flag.

```
-reduce-domains
```

Standard domain filtering affects the decision variables defined by find statements. Auxiliary variables created by SAVILE ROW are not filtered by default. To enable filtering of both find variables and auxiliary variables, use the following flag:

```
-reduce-domains-extend
```

If the problem instance contains variables with very large domains, the level of consistency is reduced from SACBounds to Minion's conventional propagation.

Aggregation of Constraints

Aggregation collects sets of constraints together to form global constraints that typically propagate better in the target solver. At present there are two aggregators for allDifferent and GCC. SAVILE ROW constructs allDifferent constraints by collecting not-equal, less-than and other shorter allDifferent constraints. GCC is aggregated by collecting atmost and atleast constraints over the same scope.

```
-aggregate
```

Mappers

The Dominion constraint solver synthesiser has mappers (also known as views) in its input language. By default SAVILE ROW will use these whenever possible. The following two flags control the use of mappers – the first switches mappers off completely, and the second allows mappers to be used only where Minion can use them. For example, multiplication-by-a-constant mappers are allowed on variables in a sum constraint, because Minion has weighted sum constraints (which, incidentally, are built from sum constraints using mappers).

```
-nomappers
-minionmappers
```

Variable Types

When creating Minion output SAVILE ROW will by default use the DISCRETE or BOOL variable type when the variable has fewer than 10,000 values and BOUND otherwise. DISCRETE and BOOL represent the entire domain and BOUND only stores the upper and lower bound, thus propagation may be lost when using BOUND. The following flag causes SAVILE ROW to use DISCRETE variables only.

```
-no-bound-vars
```

2.5.6 Controlling SAVILE ROW

The following flag specifies a time limit in milliseconds. SAVILE ROW will stop when the time limit is reached, unless it has completed tailoring and is running a solver to find a solution. The time measured is wallclock time not CPU time.

```
-timelimit <time>
```

If SAVILE ROW runs Minion it may be useful to add the following flag to limit Minion's run time, where <time> is in seconds.

```
-solver-options "-cpulimit <time>"
```

A similar approach can be used to apply a time limit to Gecode and SAT solvers.

2.5.7 Solver Control

The following flag causes SAVILE ROW to run a solver and parse the solutions produced by it. This is currently implemented for Minion, Gecode and SAT backends.

```
-run-solver
```

The following two flags control the number of solutions. The first causes the solver to search for all solutions (and SAVILE ROW to parse all solutions). The second specifies a required number of solutions.

```
-all-solutions
-num-solutions <n>
```

When parsing solutions the default behaviour is to create one file for each solution. As an alternative, the following flag will send solutions to standard out, separated by lines of 10 minus signs.

```
-solutions-to-stdout
```

The following flag passes through command-line options to the solver. The string would normally be quoted.

```
-solver-options <string>
```

For example when using Minion `-solver-options "-cpulimit <time>"` may be used to impose a time limit, or when using Gecode `-solver-options "-p 8"` causes Gecode to parallelise to 8 cores.

2.5.8 Solver Control – Minion

The following flag specifies where the Minion executable is. The default value is `minion`.

```
-minion-bin <filename>
```

When Minion is run directly by SAVILE ROW, preprocessing is usually applied before search starts. The following flag allows the level of preprocessing to be specified.

```
-preprocess LEVEL
    where LEVEL is one of None, GAC, SAC, SSAC, SACBounds, SSACBounds
```

The default level of preprocessing is SACBounds unless there are large variable domains, in which case the default is GAC.

The `-preprocess` flag affects the behaviour of Minion in two cases: first when Minion is called to filter domains (the `-reduce-domains` option), and second when Minion is called to search for a solution.

2.5.9 Solver Control – Gecode

The following flag specifies where the Gecode executable is. The default value is `fzn-gecode`.

```
-gecode-bin <filename>
```

2.5.10 Solver Control – SAT Solvers

SAVILE ROW is able to run and parse the output of MiniSAT and Lingeling solvers. The following flag specifies which family of solvers is used. Lingeling is the default.

```
-sat-family [minisat | lingeling]
```

The following flag specifies where the SAT solver executable is. The default value is `lingeling` or `minisat`, depending on the SAT family.

```
-satsolver-bin <filename>
```

In some cases the SAT output is very large and this can be inconvenient. The following option allows the number of clauses to be limited. If the specified number of clauses is reached, SAVILE ROW will delete the partial SAT file and exit.

```
-cnflimit <numclauses>
```

2.5.11 Parsing Solutions

Finally, when using ReadSolution mode, SAVILE ROW will read a solution in Minion format and translate it back to ESSENCE′. This allows the user to run Minion separately from SAVILE ROW but still retrieve the solution in ESSENCE′ format. When running Minion, the `-solsout` flag should be used to retrieve the solution(s) in a table format.

When using ReadSolution mode, the name of the aux file previously generated by SAVILE ROW needs to be specified using the `-out-aux` flag, so that SAVILE ROW can read its symbol table. Also the name of the solution table file from Minion is specified using `-minion-sol-file`. `-out-solution` is required to specify where to write the solutions. `-all-solutions` and `-num-solutions <n>` are optional in ReadSolution mode. Below is a typical command.

```
savilerow -mode ReadSolution -out-aux <filename> -minion-sol-file <filename>
                             -out-solution <filename>
```

When neither `-all-solutions` nor `-num-solutions <n>` is given, SAVILE ROW parses the last solution in the Minion solutions file. For optimisation problems, the last solution will be the optimal or closest to optimal.

Appendices

A   SAT Encoding

We have used standard encodings from the literature such as the order encoding for sums [5] and support encoding [3] for binary constraints. Also we do not attempt to encode all constraints in the language: several constraint types are decomposed before encoding to SAT.

A.1   Encoding of CSP variables

The encoding of a CSP variable provides SAT literals for facts about the variable: $[x = a]$, $[x \neq a]$, $[x \leq a]$ and $[x > a]$ for a CSP variable x and value a. CSP variables are encoded in one of three ways. If the variable has only two values, it is represented with a single SAT variable. All the above facts (for both values) map to the SAT variable, its negation, *true* or *false*. If the CSP variable is contained in only sums, then only the order literals $[x \leq a]$ and $[x > a]$ are required. Using the language of the *ladder* encoding of Gent et al [4], we have only the ladder variables and the clauses in Gent et al formula (2). Otherwise we use the full ladder encoding with the clauses in formulas (1), (2) and (3) of Gent et al. Also, for the maximum value $\max(D(x))$ the facts $[x \neq \max(D(x))]$ and $[x < \max(D(x))]$ are equivalent so one SAT variable is saved. Finally, a variable may have gaps in its domain. Suppose variable x has domain $D(x) = \{1 \ldots 3, 8 \ldots 10\}$. SAT variables are created only for the 6 values in the domain. Facts containing values $\{4 \ldots 7\}$ are mapped appropriately (for example $[x \leq 5]$ is mapped to $[x \leq 3]$). The encoding has $2|D(x)| - 1$ SAT variables.

A.2 Decomposition

The first step is decomposition of the constraints AllDifferent, GCC, Atmost and Atleast. All are decomposed into sums of equalities and we have one sum for each relevant domain value. For example to decompose AllDifferent($[x, y, z]$), for each domain value a we have $(x = a) + (y = a) + (z = a) \leq 1$. These decompositions are done after AC-CSE if AC-CSE is enabled (because the large number of sums generated hinders the AC-CSE algorithm) and before Identical and Active CSE.

The second step is decomposition of lexicographic ordering constraints. We use the decomposition of Frisch et al [2] (Sec.4) with implication rewritten as disjunction, thus the conjunctions of equalities in Frisch et al become disjunctions of disequalities. This decomposition is also done after AC-CSE and before Identical and Active CSE. However, if AC-CSE is switched on, we (independently) apply AC-CSE to the decomposition, thus extracting common sets of disequalities from the disjunctions.

The third step occurs after all flattening is completed. The constraints min, max and element are decomposed. For $\min(V) = z$ we have $V[1] = z \vee V[2] = z \ldots$ and $z \leq V[1] \wedge z \leq V[2] \ldots$. Max is similar to min with $\leq$ replaced by $\geq$. The constraint $\mathrm{element}(V, x) = z$ becomes $\forall i : (x \neq i \vee V[i] = z)$.

A.3 Encoding of Constraints

Now we turn to encoding of constraints. Some simple expressions such as $x = a$, $x \leq a$ and $\neg x$ (for CSP variable x and value a) may be represented with a single SAT literal. We have introduced a new expression type named SATLiteral. Each expression that can be represented as a single literal is replaced with a SATLiteral in a final rewriting pass before encoding constraints. SATLiterals behave like boolean variables hence they can be transparently included in any constraint expression that takes a boolean subexpression.

For sums we use the order encoding [5] and to improve scalability sums are broken down into pieces with at most three variables. Sum-equal constraints are split into sum-$\leq$ and sum-$\geq$ before encoding. For other constraints we used the standard support encoding wherever possible [3]. Binary constraints such as $|x| = y$ use the support encoding, and ternary functional constraints $x \diamond y = z$ (eg $x \times y = z$) use the support encoding when z is a constant. Otherwise, $x \diamond y = z$ are encoded as a set of ternary SAT clauses: $\forall i \in D(x), \forall j \in D(y) : (x \neq i \vee y \neq j \vee z = i \diamond j)$. When $i \diamond j$ is not in the domain of z, the literal $z = i \diamond j$ will be false. Logical connectives such as $\wedge, \vee, \leftrightarrow$ have custom encodings and table constraints use Bacchus' encoding [1] (Sec.2.1).

References

[1] Fahiem Bacchus. Gac via unit propagation. In *Principles and Practice of Constraint Programming–CP 2007*, pages 133–147. Springer, 2007.

[2] A.M. Frisch, B. Hnich, Z. Kiziltan, I. Miguel, and T. Walsh. Global constraints for lexicographic orderings. In P. van Hentenryck, editor, *Proceedings of the Eighth International Conference on Principles and Practice of Constraint Programming*, pages 93–108, 2002.

[3] Ian P. Gent. Arc consistency in SAT. In *Proceedings of the 15th European Conference on Artificial Intelligence (ECAI 2002)*, pages 121–125, 2002.

[4] Ian P. Gent and Peter Nightingale. A new encoding of alldifferent into sat. In *Proc. 3rd International Workshop on Modelling and Reformulating Constraint Satisfaction Problems (ModRef 2004)*, pages 95–110, 2004.

[5] Naoyuki Tamura, Akiko Taga, Satoshi Kitagawa, and Mutsunori Banbara. Compiling finite linear CSP into SAT. *Constraints*, 14(2):254–272, 2009.